


Qiwei Chen

SOFTWARE ENGINEERING UNDERGRADUATE · EMBODIED AI & ML SYSTEMS

Shanghai Jiao Tong University, Shanghai, China

☎ (+86) 131-2312-8852 | ✉ qiweic10@sjtu.edu.cn | 🌐 kiwi142857 | 🏠 Homepage



Education

Shanghai Jiao Tong University

Shanghai, China

B.S. IN SOFTWARE ENGINEERING

Aug. 2022 - Jun. 2026

- Academic score: 91.1/100 (rank 8/89); GPA: 3.92/4.30 (rank 9/89).
- Honors include Jiangshi Scholarship, 85th CS Education Development Fund & Yang Yuanqing Education Fund Scholarship, Gold Medal at China International College Students' Innovation Competition (Shanghai Division), Huawei RoboMind Winter Camp First Prize, and OpenHarmony Training Camp National Third Prize.

Publication

Characterizing Vision-Language-Action Models across XPU: Constraints and Acceleration for On-Robot Deployment

ICML

KAIJUN ZHOU, QIWEI CHEN, DA PENG, ZHIYANG LI, XIJUN LI, AND JINYU GU

2026

Selected Projects

VLA/VLM Edge Deployment and Inference Optimization

ASCEND 310P, RK3588, OPENHARMONY, LLAMA.CPP

May 2025 - Aug. 2025

- Adapted and optimized VLA, VLM, and SAM2 models on domestic edge accelerators under tight latency and memory constraints.
- Implemented RK3588 operator and memory optimizations on OpenHarmony, and upstreamed a llama.cpp-based VLM optimization to the open-source community.

CANN Embodied Intelligence Community Contribution

VLA MODELS, ASCEND CANN, OPEN-SOURCE COMMUNITY

2025

- Implemented Ascend platform recipes for VLA models including pi0, ACT, and Diffusion Policy, improving reproducibility for embodied-AI deployment on Ascend hardware.

ICS Course Lab Optimization

TRANSFORMER PARALLELISM, KV CACHE, GRADING INFRASTRUCTURE

Aug. 2024 - May 2025

- Designed LLM Lab modules for Transformer parallelism and KV-cache optimization, with evaluation workflows for student implementations.
- Migrated the course platform from SVN to Git and built an online grading and score-query system.

OpenTrustee LLM Secure On-Device Inference

TEE, OPENTRUSTEE, LLAMA.CPP

Sep. 2024 - Nov. 2024

- Built trusted-execution model loading and inference modules to protect model parameters, then reduced bottlenecks through parallel model loading and decryption.

Minik8s Lightweight Orchestration System

TEAM LEAD; KUBERNETES-STYLE SCHEDULER AND SERVICES

Mar. 2025 - Jun. 2025

- Led development of a simplified Kubernetes cluster orchestrator with pod scheduling, service discovery, and load balancing.

Skills

Languages

C/C++, Python, Java, Bash

ML Systems

VLA models, LeRobot, llama.cpp, vLLM, Transformer optimization, edge inference

Systems

Distributed systems, Kubernetes, TEE, OpenHarmony, Ascend CANN

Languages

CET-6; strong technical reading and writing in English